

Mathematical Statistics II

Introduction to Point Estimation

Jesse Wheeler

Introduction

Overview

- We will formally introduce the idea of point estimation.
- In addition to an introduction, we will introduce the concept of the empirical distribution, as well as methods of moment estimators.
- The material for this section largely comes from Chapter 8 of Rice (2007).

Point Estimation: An introduction

Point estimation

- In the previous lecture(s), we provided an example of Bayesian vs Frequentist point-estimation via first principles.
- That is, using the various interpretations, we could reason an estimate for the probability p in a binomial experiment.
- We are now interested in studying approaches for more general cases.
- Given a dataset and a chosen model, how can we estimate parameters?

Point estimation II

- We will first start with some notation, and motivating examples.
- Term *model* in this class will generally refer to a probability model, and can be based on a discrete or continuous probability measure.

Point estimation III

Normal Model

The Normal (or Gaussian) family of distributions arises often in the real world. Examples include human heights (conditioned on gender), rainfall amounts, and many biological measurements are approximately normal (or log-normal).

Given a set of observations x_1, x_2, \dots, x_n , we may *model* these as iid normal $X_i \sim N(\mu, \sigma^2)$, and our goal being using the data to estimate the values of μ or σ .

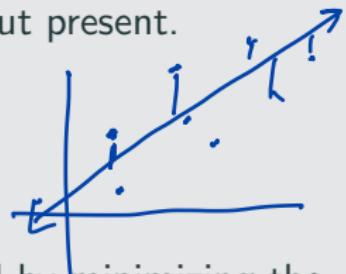
Point estimation IV

Regression

Sometimes the probability model is *implicit*, but present.

Consider the regression model:

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$$



We often think of fitting this regression model by minimizing the average squared-error: $(Y_i - \hat{Y}_i)^2$. However, this approach typically corresponds to an implicit probability model for the error terms ε_i , namely a normal distribution with mean 0. In this case, we might want to estimate β_0 , β_1 , and σ^2 , which is $\text{Var}(\varepsilon_i)$.

fit β_0, β_1 by $\min : \sum (Y_i - \hat{Y}_i)^2$
 $\Leftrightarrow Y|X \sim N(\beta_0 + \beta_1 X, \sigma^2)$

$$y_i = f(x_i) + \epsilon_i$$

$\underbrace{f(x_i)}$
NN

$$\min \sum (\hat{f}(x_i) - y_i)^2$$

$$y|x_i \sim N(f(x_i), \sigma^2)$$

$\underbrace{\quad\quad\quad}$

Point estimation V

Poisson Process

Another common example is a Poisson Process model. Many real-world phenomena are well-approximated by a Poisson process, over space or time. Examples include arrival times at a gas station, number of meteors landing in a geographic area, radioactive decay, etc. Here, there is only one parameter we want to estimate using data, namely the rate λ .

Parameter Estimation

- All of the above examples have the common feature that we pick a *model*, and we want to use the model to describe the data-generating process.
- More accurately, however, we pick a candidate **family** of models; (Gaussian family, Poisson Family, Linear Regression family, etc).
- Generally, the exact model needed within a *family* of models is determined by a few parameters.

Example: Gamma-Rainfall

- The Gamma distribution depends on two parameters, α and λ :

$$f_X(x; \alpha, \lambda) = \frac{1}{\Gamma(\alpha)} \lambda^\alpha x^{\alpha-1} e^{-\lambda x}.$$

- The Gamma distribution is quite flexible, and works as a useful model for various situations.
- One example is modeling rainfall amounts per-storm under two conditions, cloud seeding vs not cloud seeding (simulated data, couldn't find original data).
- A Gamma distribution fits both samples well, but we get different parameters α and λ for the two different samples
- Differences in the respective distributions are reflected in differences in the parameters α and λ .

Two-sample Rainfall

(1): $\text{Gamma}(\alpha_1, \gamma_1)$

(2): $\text{Gamma}(\alpha_2, \gamma_2)$

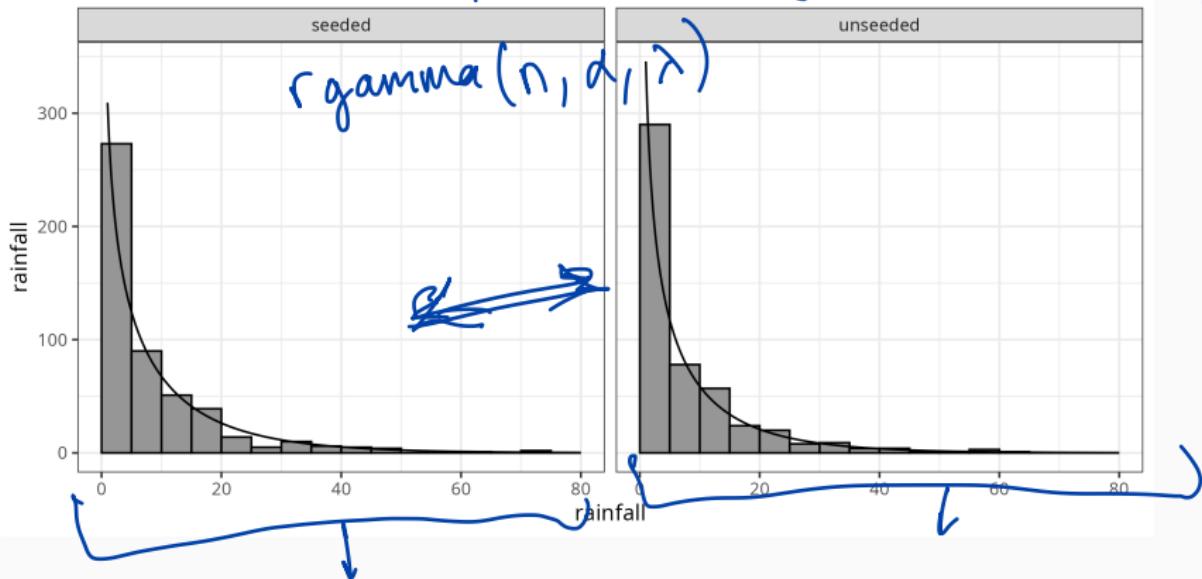


Figure 1: Data and model fit to two different Gamma distributions.

Notation and generalizations

$$X^* \in \mathbb{R}^{n \times d}$$

- We will generalize by using the following ideas and notations.
- We will denote the *observed data* as $x_1^*, x_2^*, \dots, x_N^*$, and use the shorthands $x_{1:N}^*$ if we emphasize the entire collection, and x^* if the emphasis is not needed.
- We assume that the data are realizations of random variables X_1, X_2, \dots, X_N , again using the notation $X_{1:N}$ for the collection of N random variables, or X if this is not needed.
- In general, the data x_i^* and random variables X_i can be multivariate, but focus primarily on the univariate case.

$$x_i^* \in \mathbb{R}^d$$

Notation and generalizations II

- We will be interested in fitting a probabilistic model $f_{X_{1:N}}(x_{1:N}; \theta)$ using the data. The model may correspond to a discrete probability, or a continuous probability. In these cases, f is usually a pmf or pdf, respectively.
- Subscripts will be dropped occasionally if it is not necessary. For instance, $f(x; \theta)$ is taken to mean the model of all data $x = x_{1:N}$, and would formally be expressed as $f_{X_{1:N}}(x_{1:N}; \theta)$.
- This approach is sometimes called “function overload”; it’s not my favorite approach, but it is convenient. The meaning of the function is primarily understood by the arguments and context.
- The function $f(x; \theta)$ belongs to a particular **family** of models, indexed by θ , which is generally multivariate.

Notation and generalizations III

Normal model example

Suppose we observe the following data: 3.49, 2, 3.38, 1.62, 2.18, and we would like to fit a normal model to the data, assuming the data are iid. Then $x_1^* = 3.49$, $x_2^* = 2$, and so forth, and the model family depends on $\theta = (\mu, \sigma^2)$, and the model can be expressed as:

$$f(\mathbf{x}; \theta) = f_{X_{1:N}}(x_{1:N}; \theta)$$

$$\begin{aligned} f(x; \theta) &= f_{X_{1:5}}(x_{1:5}; \mu, \sigma^2) \\ &= \prod_{i=1}^5 f_{X_i}(x_i; \mu, \sigma^2) \quad X_i \rightarrow x_i \\ &= \prod_{i=1}^5 \frac{1}{\sigma \sqrt{2\pi}} e^{-(x_i - \mu)^2 / 2\sigma^2} \end{aligned}$$

Our goal is to estimate μ , σ^2 using the observed data $(x_{1:5}^*)$.

Notation and Generalization (continued)

- Our goal now is to develop general procedures for estimating θ , using observed data x^* , and a proposed family of models $f(x; \theta)$.
- We will develop three main approaches: (1) Method of Moments (2) Maximum Likelihood Estimation, and (3) Bayesian estimation.
- In this section, we will focus only on method of moments estimators.
- Once point estimation techniques are developed, we will provide theory about these estimates and their uncertainty; discussing bias, variance, and optimality of estimates.

Brief Introduction to R

R introduction

- Before we start looking at real-data examples, let's introduce some basic R coding principles that will help us calculate moments from the data.
- R is a programming language, but for the sake of this class, we'll just treat it as a statistics calculator.
- For now, we will only focus on the most simple data types and operations: creating objects, vectors, and computing summary statistics.

R introduction II

- First, saving objects in R. We can use = (like most languages), or the assignment operator: <-

```
x <- 2
```

```
x + 2
```

```
[1] 4
```

R introduction III

- A vector in R is a collection of objects of the same data type.
In this class, we will only need to use numeric data types

R introduction IV

```
x <- c(1, 2, 3, 4, 5)  
class(x)
```

```
[1] "numeric"
```

```
mean(x)
```

```
[1] 3
```

```
sum(x)
```

```
[1] 15
```

R introduction V

- Some fast ways of building vectors include:

```
1:5 # this gives 1, 2, 3, 4, 5
```

```
[1] 1 2 3 4 5
```

```
seq(1, 10, by = 2) # Gives 1, 3, 5, 7, 9
```

```
[1] 1 3 5 7 9
```

seq(1, 10, length.out = 1000)

R introduction VI

- For generating random numbers, we can use the syntax:
`rdist.`

`r pois(n=10, rate=5)`

↓
10 obs from `pois(5)`

`r norm(n, mean, sd)`

? `r beta`

`r gamma . . .`

`r beta . . .`

`r unif . . .`

R introduction VII

```
rnorm(n = 10, mean = 2, sd = 1)
```

```
[1] 1.7531041 0.7844391 3.5614051 2.4273102 0.7989765 3.0524585  
[7] 0.6949364 1.3073924 2.6026489 1.8022469
```

```
rpois(n = 7, lambda = 5)
```

```
[1] 2 6 1 5 5 9 6
```

```
rbeta(n = 3, shape1 = 0.8, shape2 = 1.3)
```

```
[1] 0.51652672 0.10386537 0.05986089
```

R introduction VIII

- Lastly (and maybe most important), function documentation and help is readily available by appending a question mark:
`?rnorm`

```
?mean
```

```
?rnorm
```

```
?sd
```

Method of Moments

Motivation

- The Method of Moments (MoM) estimation technique is a simple idea.
- Pick a family of models $f(x; \theta)$, and observed data x^* . $X \sim \text{pois}(x)$
- The family of models will have theoretical moments, i.e., $E[X^k]$.
- Generally, these moments can be expressed in terms of the model parameters, θ .
- Thus, we will estimate $\hat{\theta}$ so that the data moments match the theoretical moments.

$$E[X] = \textcircled{X}$$

$$X \sim N(\mu, \sigma^2)$$

$$\rightarrow E[X] = \textcircled{\mu}$$
$$\rightarrow E[X^2] = \textcircled{\mu^2 - \sigma^2}$$

The empirical distribution

- One justification of this approach considers the **empirical distribution** of observed data.
- Let X_1, X_2, \dots, X_N be random variables, representing a possible data sample.
- We will assume that X_i are iid, from some distribution F_θ (F_θ is the cdf here).
- We will define the empirical distribution function as:

$$F_N(t) = \frac{1}{N} \sum_{i=1}^N I[X_i \leq t].$$

$\frac{1}{N}$ # of obs smaller than t .

- When we observe a specific dataset x^* , we can plug in these numbers to get a specific distribution that is not random.

The empirical distribution II

- A few things to note is that $F_N(t)$ is a proper CDF.
- By the law of large numbers, $F_N(t) \xrightarrow{a.s.} F_\theta(t)$ for every point t .
- The Glivenko–Cantelli theorem also strengthens this statement by saying that the convergence is uniform, in the sense that $\sup_t |F_n(t) - F_\theta(t)|$ converges to zero.

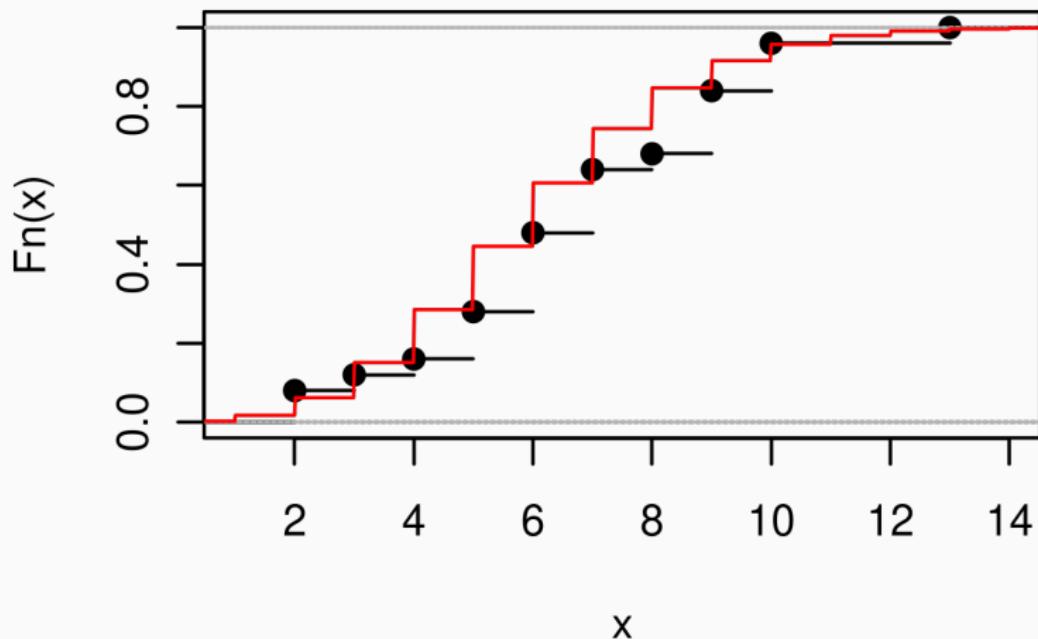
The empirical distribution III

Example: empirical distribution function for Poisson Data:

```
set.seed(123) # Reproducible results
x <- rpois(n = 25, lambda = 6)
plot(ecdf(x))
lines(
  x = seq(1e-16, 15, length.out = 1000),
  y = ppois(seq(1e-16, 15, length.out = 1000), 6),
  col = 'red'
)
```

The empirical distribution IV

$\text{ecdf}(x)$



Method of Moments Estimation

- It can be shown that the k th moment of the empirical distribution is

$$\hat{\mu}_k = \frac{1}{N} \sum_{i=1}^N X_i^k. \rightarrow E_{\theta} [X_i^k]$$

- Method of Moments idea:

- For many commonly used parametric families (e.g., Gaussian, Poisson), the distribution is completely specified by a small set of parameters.
- These parameters are typically explicit functions of the moments of the distribution (e.g., mean and variance for the Gaussian).

Method of Moments Estimation II

- Although the moment generating function (MGF) uniquely determines the entire distribution, in many model families, the relevant parameters are uniquely determined by just the first few moments.
- Therefore, as the empirical moments computed from data converge to the true moments (by the Law of Large Numbers), it is natural to estimate model parameters by equating empirical and theoretical moments—leading to the method of moments estimators.

Method of Moments: generalized version

- To summarize mathematically, let $\mu_k = E[X^k]$ be the theoretical k th moment.
- Let $\hat{\mu}_k = \frac{1}{N} \sum_{i=1}^N X_i^k$ be the k th sample moment.
- $\hat{\mu}_k$ is an estimate of μ_k ; however, we don't want an estimate of μ , we want an estimate of θ !
- For models with finite parameters, $\theta = (\theta_1, \dots, \theta_k)$, we can often express θ_i as a function of (μ_1, \dots, μ_k) :

$$\theta_i = g_i(\mu_1, \dots, \mu_k)$$

.

Method of Moments: generalized version II

- Thus, our estimate of θ_i would be found by plugging in the empirical moments:

$$\hat{\theta}_i = g_i(\hat{\mu}_1, \dots, \hat{\mu}_k).$$

Examples

Poisson Distribution

Suppose we observe data $x_{1:N}^*$, and want to fit a Poisson model. Since the Poisson distribution only has one parameter (λ), our goal is to use x^* to estimate λ .

$$\theta = (\theta_1) \quad E[x_i] = \underline{\lambda} = \theta_1$$

$$g_1(\underline{\mu}_1) = g_1(E[x_i]) = \lambda$$

$$g_1(\hat{\mu}_1) = \left(\frac{1}{N} \sum_i x_i \right) = \hat{x}$$

$$\hat{\lambda}_{MOM} = \bar{x}_N$$

Real-data example

Poisson distribution with real data

The National Institute of Science and Technology collected data about asbestos fibers on filters. Asbestos dissolved in water was spread on a filter, and the number of fibers in each of 23 grid squares were counted:

```
[1] 31 29 19 18 31 28 34 27 34 30 16 18 26 27 27 18 24 22  
[19] 28 24 21 17 24
```

$$\hat{\lambda} = \frac{1}{N} \sum x_i = \bar{x}_N$$

Real-data example II

The mean can be calculated as:

$$\text{mean}(x) = \hat{\lambda} \approx 24.9$$

[1] 24.91304

- What about the error associated with this estimate?

$$x_i \sim \text{Pois}(\lambda)$$

$$\hat{\lambda} = \frac{1}{N} \sum x_i$$

$$\sum_{i=1}^N x_i \sim \text{Pois}(N\lambda)$$

$$N\hat{\lambda} \sim \text{Pois}(N\lambda)$$

Sampling Distribution

$$N\hat{x} \sim \text{Pois}(N\lambda)$$

- As always, we are interested in the uncertainty related to our estimates.
- In most cases, we cannot directly calculate uncertainty of estimates, and we will have to rely on more advance theory, discussed later.
- Sometimes, however, we can calculate some form of uncertainty based on the form of the estimator, and model assumptions.
- The last (and next) models are such cases.

$$\text{E}(\hat{x}) = \text{Var}(\hat{x})$$

$$\stackrel{?}{=} \text{SD}(\hat{x})$$

$$\text{Var}(N\hat{x}) = N\lambda$$

$$\text{Var}(\hat{x}) = \frac{N\lambda}{N^2}$$



Sampling Distribution II

Sampling Distribution

Most estimates $\hat{\theta}$ of θ are functions of the random variables X_1, X_2, \dots, X_N . Thus, $\hat{\theta}$ is also a random variable. The distribution of $\hat{\theta}$ is called the **sampling distribution**.

- In most cases, the exact distribution of $\hat{\theta}$ is unknowable.
- Instead, we often get approximations to this distribution, and in particular, the variance of the distribution, in order to quantify uncertainty of the estimator.
- For the Poisson model and the method of moments estimator, however, we can calculate this exactly.

Sampling Distribution III

Sampling distribution of Poisson MoM estimator

Let X_1, \dots, X_N be modeled as iid from a $\text{Poisson}(\lambda)$ distribution. If $\hat{\lambda}$ is the method of moments estimator of λ , what is its sampling distribution?

Sampling Distribution IV

For our specific dataset, we can approximate the standard error as:

```
sqrt(mean(x) / length(x)) |> round(2)
```

```
[1] 1.04
```

The “pipe” operator in R `|>` takes the output of one function, and inputs it as the first argument into the next function.

Example: Model Checking

- In addition to estimating the stochastic uncertainty due to sampling distribution, we might want to assess inductive uncertainty due to model selection.
- This is a large topic, one we will revisit in much more detail later.
- For now, I want to present a simple idea that can be done with what we already know.
- This will be based on a **parametric bootstrap**, a technique we will discuss in detail later in the course.

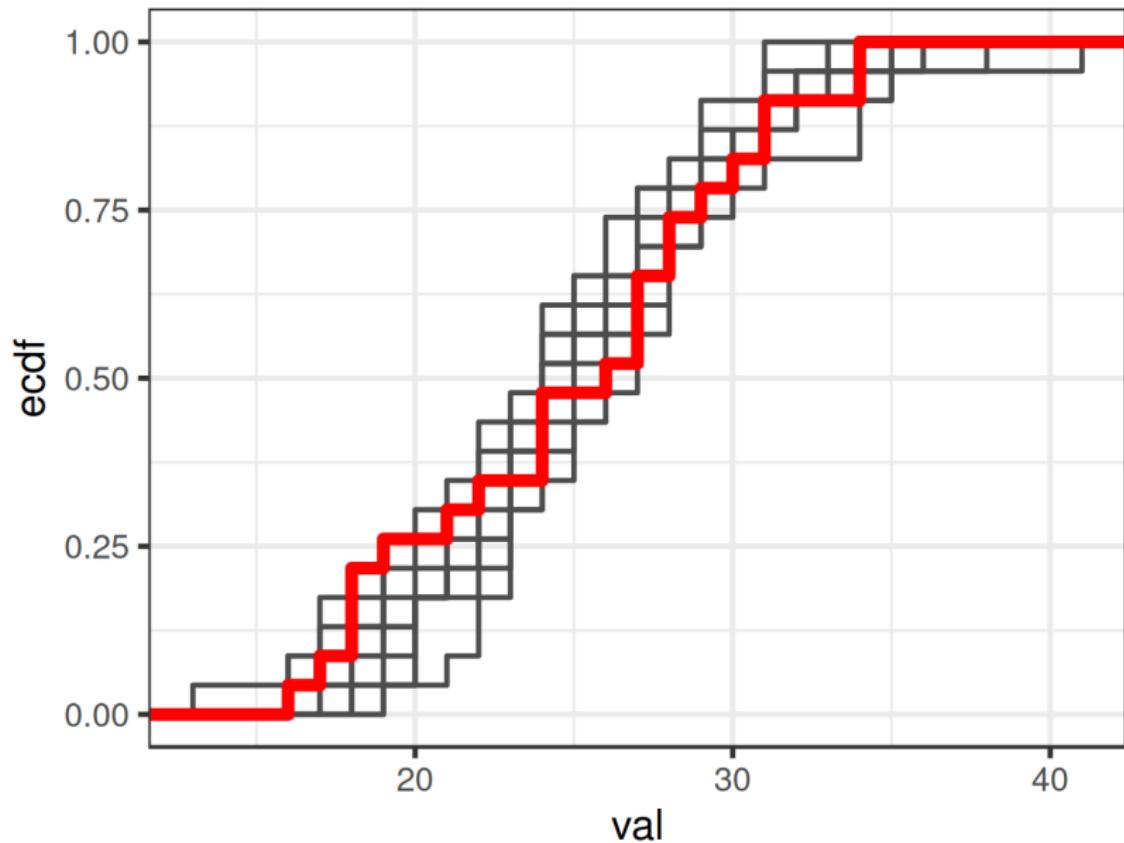
Example: Model Checking II

- The idea: we claim that the data come from a $\text{Poisson}(\lambda)$, and we estimated $\hat{\lambda}$.
- If the data really are $\text{Poisson}(\lambda)$, then our estimate $\hat{\lambda} \approx \lambda$.
- We can simulate many different data from $\text{Poisson}(\hat{\lambda})$, and compare this to our real data.
- If the model is reasonable, then our data shouldn't look too different from the simulations.
- We can compare the empirical distribution functions from the real data, and simulations.

ECDF: R code

```
library(tidyverse)
lambda <- mean(x)
results <- replicate(10, rpois(length(x), lambda)) |>
  as.data.frame()
colnames(results) <- paste0("sample_", 1:10)
results |>
  pivot_longer(
    cols = everything(), names_to = "replicate",
    names_prefix = "sample_", values_to = "val"
  ) %>%
  ggplot(aes(val, group = replicate)) +
  stat_ecdf(geom = "step", col = 'grey30') +
  stat_ecdf(
    data = data.frame(val = x, replicate = '0'),
    geom = 'step', col = 'red', linewidth = 1.2
  ) + theme_bw()
```

Plot output



Example: Normal Distribution

Normal Distribution

Suppose we observe N observations, and we want to model them as iid $N(\mu, \sigma^2)$. Find the method of moments estimator $\theta = (\mu, \sigma^2)$.

Example: Gamma Distribution

MoM estimator for Gamma distribution

Consider modeling the data X_1, \dots, X_N as iid $\text{Gamma}(\alpha, \lambda)$.

Find the method of moments estimator for these data.

Example: Muon Decay

MoM estimator for Muon Decay

In *statistical physics*, the angle ϕ at which electrons are emitted in muon decay is modeled using the following density:

$$f(x; \alpha) = \frac{1 + \alpha x}{2}, \quad -1 \leq x \leq 1,$$

and where the parameter α satisfies $-1 \leq \alpha \leq 1$, and $x = \cos \phi$. Supposing we observe data X_1, X_2, \dots, X_N , what is the method of moments estimator for α ?

References and Acknowledgements

Rice JA (2007). *Mathematical statistics and data analysis*, volume 371. 3 edition. Thomson/Brooks/Cole Belmont, CA.

- Compiled on January 15, 2026 using R version 4.5.2.
- Licensed under the [Creative Commons Attribution-NonCommercial license](#). Please share and remix non-commercially, mentioning its origin.
- We acknowledge [students and instructors for previous versions of this course / slides](#).

