Mathematical Statistics II

Maximum Likelihood Estimation

Jesse Wheeler

Introduction

Overview

- The next approach we will discuss is Maximum Likelihood Estimation (MLE).

- As we will see, the MLE has several desirable properties, and as a result is often favored over approaches like the method of moments.

- The material for this section largely comes from Chapter 8.5 of Rice (2007), and various sections in Pawitan (2001).

Likelihood: an introduciton

What is likelihood?

- The term "likelihood" is often used colloquially to mean something analogous to probability. E.g., "What is the likelihood that it rains tomorrow?"

- When we use this term in statistics / mathematics, we mean something specific that isn't the same thing as probability.

- The use of the term "likelihood" was first made by R. A. Fisher, who was the architect and primary proponent of "likelihood-based-inference".

- We will start with the treatment of likelihood in the text "In all Likelihood" (Pawitan, 2001), which is a fantastic resource on the subject. (This will lead to some review...)

Coin Flips

We will revisit this example, as it is a great starting point to connect with existing understanding.

Consider flipping a coin $N = 10$ times.

What is likelihood? III

- For our specific coin-flipping example with $N = 10$, $X = 8$, the likelihood function is

$$L(\theta) = P_\theta(X = 8).$$

- This is plotted in the following way:

```
x <- seq(1e-8, 1-1e-8, length.out = 1000)
y <- dbinom(8, 10, x)
plot(x = x, y = y, type = 'l')
```

What is likelihood? V

- From the figure, we see that p is unlikely to be less than 0.5,or greater than 0.95.

- Given the data alone (no prior), we should prefer a value somewhere in the middle of these values.

- We still have some uncertainty about the value of p, but the likelihood gives us a numerical way to compare values of θ. Stochastic uncertainty as a result of sampling is captured in the likelihood function $L(\theta)$.

What is likelihood? VI

- The likelihood is not a probability. Though it came from a probability, the likelihood function (a function of θ) does not satisfy the requirements to be a probability. In our previous example, we have:

$$\int_0^1 L(\theta)\, d\theta = 1/11 \neq 1.$$

- For discrete probability, the likelihood was continuous. Discrete likelihoods are possible, arising when we want to select from a list $\{\theta_1, \theta_2, \ldots\}$.

What is likelihood? VII

- The idea behind maximum likelihood estimation (MLE) is simple: our estimate is the value of θ that maximizes the likelihood function $L(\theta)$.

- The MLE is considered a *frequentist* approach. Why? It quantifies a maximum belief about a parameter, which is more Bayesian in nature than Frequentist.

- As we'll see later, the MLE has nice theoretical Frequentist *properties*, and as a result can be justified via the frequentist paradigm.

- Still, it has close connection to Bayesian estimation and interpretation. In fact, we'll discuss connections between the MLE and Bayesian statistics later.

What is likelihood? VIII

- Often, maximizing the likelihood directly is challenging, so we maximize the log-likelihood instead.

- Other times, the likelihood has to be maximized numerically.

MLE of coin toss problem

Suppose we have $N = 10$ total tosses, and n total heads. Find the MLE of p, the probability of heads.

Continuous models

- The interpretation of the likelihood function as the "the probability of the observed data x^*, considered as a function of θ" makes perfect sense in the discrete model case.

- For continuous models, the technical issue arises that the probability of any point value x is zero.

- We resolve the problem similar to what was done in Math 4450 and the John Rice text: approximate the probability by discretizing into small, discrete intervals:

$$x^* \in (x^* - \epsilon/2, x^* + \epsilon/2),$$

Continuous models II

thus, the probability of observing something ϵ-close to the data is:

$$L(\theta) = P_\theta\big(X \in (x^* - \epsilon/2, x^* + \epsilon/2)\big)$$
$$= \int_{x^* - \epsilon/2}^{x^* + \epsilon/2} f(x; \theta) \, d\theta \approx \epsilon f(x^*; \theta).$$

- Then, since the likelihood is only meaningful up to a constant (we will discuss likelihood ratios later), then this has the same behavior as $L(\theta) = f(x^*; \theta)$.

- There are more advanced approaches to this problem, but this simple argument justifies the use of the pdf of a continuous random variable as the likelihood $L(\theta)$.

Continuous models III

- **Going forward:** we once again will generalize a model $f(x; \theta)$ to mean either the pmf or pdf of a random variable. I will often say "density" as a blanket term, even if this corresponds to a pmf, not a density.

- Further, when we "integrate" a density, this means either:

$$\int f(x; \theta)\, dx, \quad \text{If continuous}$$

or

$$\sum_x f(x; \theta), \quad \text{If discrete.}$$

Joint Probabilities

Likelihood with multiple observations

- Often the data we observe is multi-dimensional, rather than summarized as a single observation.

- In this case, the likelihood θ is still determined via the joint model:

$$L(\theta) = f(x^*; \theta) = f_{X_{1:N}}(x_1^*, x_2^*, \ldots, x_N^*; \theta).$$

- We are mostly focused in this class in the case were the observations are independent, meaning the likelihood factors:

$$L(\theta) = \prod_{i=1}^{N} f_{X_i}(x_i^*; \theta).$$

Likelihood with multiple observations II

- We often further simplify this by assuming the data are identically distributed:

$$L(\theta) = \prod_{i=1}^{N} f_{X_1}(x_i^*; \theta).$$

- As we've seen, it's generally easier to maximize the log-likelihood. In the IID case:

$$\ell(\theta) = \log \prod_{i=1}^{N} f_{X_1}(x_i^*; \theta) = \sum_{i=1}^{n} \log f_{X_1}(x_i^*; \theta).$$

Examples

Examples of finding the MLE

Traffic data: Poisson Model

Returning to a motivating example, suppose we model traffic accidents in a given week as $X_1, X_2, \ldots, X_N$, where the data are iid Poisson(λ). Obtain the MLE for λ.

Examples of finding the MLE II

Two parameter model: Gaussian model

Suppose we model observations $X_1, \ldots, X_N$ as IID $N(\mu, \sigma^2)$ random variables. Find the MLE of $\theta = (\mu, \sigma^2)$.

Plotting Normal Likelihood

- The likelihood function (not just to point estimate) will be used to measure uncertainty.
- For models with a single parameter, we often plot the likelihood curve.
- With more than one parameter, however, we have a likelihood surface.
- For the iid Normal(μ, σ^2) model, code for plotting this surface is available with course source-code.

Figure 1: Likelihood surface of data generated from normal distribution.

Calculating the likelihood in R

- For standard distributions, the likelihood is easy to calculate in R. Recall the likelihood is:

$$L(\theta) = f(x_i^*; \theta).$$

- For $X_1, \ldots, X_n$ iid normal, the likelihood can be calculated using dnorm:

```r
# Synthetic data
x_star <- rnorm(n = 100, mean = 4, sd = 3)
theta <- c(0, 1)  # value of theta
Likelihood <- prod(dnorm(x, mean = theta[1], sd = theta[2]))
loglik <- sum(
  dnorm(x, mean = theta[1], sd = theta[2], log = TRUE)
)
```

Numeric Optimization

Numeric Optimization

- In the previous examples, the MLE was available *analytically*.

- In many cases, however, there is no closed-form solution for the MLE, and it must be computed numerically.

- The next example demonstrates this, and then we will discuss optimization strategies.

Numeric Optimization II

Example: Gamma likelihood

Suppose we want to model data $X_1, X_2, \ldots, X_n$ as iid
Gamma(α, λ), which has the density function:

$$f(x;\, \alpha, \lambda) = \frac{1}{\Gamma(\alpha)} \lambda^\alpha x^{\alpha-1} e^{-\lambda x}, \quad 0 \le x < \infty.$$

Find the MLE of $\theta = (\alpha, \lambda)$.

Numeric Optimization III

- The previous example leads us to consider numeric techniques for optimization and root finding.

- Note that this class is not an optimization course, so we'll only cover some of the most basic ideas.

- More modern and efficient numeric optimization techniques are readily available in R (or any other statistical software).

- For this class, we'll introduce some basic ideas like the Newton-Raphsom approach for root finding and optimization, as well as some other basic methods.

Newton-Raphsom root-finding algorithm

- Idea: start at a point θ_0, and approximate find the tangent line of $f(\theta)$ at the point θ_0:

$$y - f(\theta_0) = f'(\theta_0)(\theta - \theta_0)$$

- Then, find the root of the tangent line by setting $y = 0$, and solving for θ:

$$\theta = \theta_0 - \frac{f(\theta_0)}{f'(\theta_0)}.$$

Newton-Raphsom root-finding algorithm II

- This root of the tangent line will be closer than our original guess θ_0, so we set:

$$\theta_1 = \theta_0 - \frac{f(\theta_0)}{f'(\theta_0)},$$

and repeat:

$$\theta_{n+1} = \theta_n - \frac{f(\theta_n)}{f'(\theta_n)}.$$

- We stop based on some convergence criteria, often something like $|\theta_{n+1} - \theta_n| < \epsilon$, for a small choice of ϵ.

- (In class, check out wikipedia or draw a picture).

Newton-Raphsom root-finding algorithm III

- We need now a starting point θ_0.

- Really we can pick anything, but it's best if we are close to the MLE.

- For our current problem (Gamma distribution), we could use the MoM estimator:

$$\hat{\alpha}_{\mathsf{MoM}} = \theta_0 = \frac{\left(\bar{x}_n^*\right)^2}{\frac{1}{n}\sum_{i=1}^n \left(x_i^* - \bar{x}_n^*\right)^2}.$$

Newton-Raphsom root-finding algorithm IV

```r
NR_root <- function(theta0, fn, deriv, tol = 1e-8, maxiter = 100
  iter <- 0
  theta_old <- theta0
  theta_new <- theta_old + 10 * tol

  while(abs(theta_old - theta_new) > tol && iter < maxiter) {
    iter <- iter + 1
    theta_old <- theta_new
    theta_new <- theta_old - fn(theta_old) / deriv(theta_new)
  }
  cat("iters: ", iter, "\n")
  theta_new
}
```

Newton-Raphsom root-finding algorithm V

```
alpha_fn <- function(alpha, data) {
  n <- length(data)
  n * log(alpha) - n * log(mean(data)) + sum(log(data)) - n * di
}

alpha_deriv <- function(alpha, data) {
  n <- length(data)
  (n/alpha) - n * psigamma(alpha, 1)
}

set.seed(123)
data <- rgamma(n = 23, shape = 1, rate = 2)
```

Newton-Raphsom root-finding algorithm VI

```r
# Not the exact MoM estimate, but close enough:
alpha_mom <- (mean(data)^2) / sd(data)

NR_root(
  theta0 = alpha_mom,
  fn = function(x) alpha_fn(x, data = data),
  deriv = function(x) alpha_deriv(x, data = data),
  tol = 1e-10
)

iters:  7
[1] 0.9728019
```

Newton-Raphsom root-finding algorithm VII

- Once the estimate of α is found, we can then plug it in to get the estimate of λ:

$$\hat{\lambda} = \frac{\hat{\alpha}}{\bar{x}_n^*} = 1.981.$$

Figure 2: Likelihood surface of data generated from Gamma distribution.

Root-find considerations

- The function that we built for solving $f(\theta) = 0$ requires the derivative $f'(\theta)$.

- In some cases, the derivative is not readily available. Instead, we can approximate using the definition:

$$f'(\theta) = \lim_{h \to 0} \frac{f(\theta + h) - f(\theta)}{h} \approx \frac{f(\theta + \Delta) - f(\theta)}{\Delta}.$$

- We just pick Δ to be small, and we get a very good approximation of the derivative.

- Thus, we don't really need to derivative for uni-variate root-finding.

- The same mathematical approach can readily be extended into higher dimensional θ, replacing the derivative with a gradient.

Root-find considerations II

- An important consideration, however, is that the results may depend on your starting parameter θ_0. In practice, you may want to try multiple values of θ_0.

Root-find considerations III

- In most programming languages, there will be pre-built methods for root-fining and optimization.

- Our function we built works, but it doesn't do careful error checking, and it isn't optimized for speed.

- For univariate-root finding $f(\theta) = 0$ in R, we can use the uniroot function, which doesn't require a derivative.

- This function is very efficient for solving roots if $\theta \in \mathbb{R}$. For higher dimensions, we need to import a different package.

Root-find considerations IV

```
uniroot(
  function(x) alpha_fn(x, data = data),
  interval = c(0.5, 2)
)

$root
[1] 0.9728068

$f.root
[1] -7.754612e-05

$iter
[1] 6

$init.it
[1] NA
```

Direct Numeric Optimization

- Based on our last example, we solved for one parameter λ, and numerically found the other, α, via root-finding.
- In many cases, this is not possible. Instead, we might want to directly optimize both parameters.
- We will first introduce this by extending the Newton-Raphson to perform optimization rather than root-finding.

Direct Numeric Optimization II

- The idea is simple: local maximum are just the zeros (roots) of the derivative function.

- Thus, we still perform Newton-Raphson, but on the derivative rather than the original function.

- Starting with initial estimate θ_0, we update following the equations until convergence:

$$\theta_{n+1} = \theta_n - \frac{f'(\theta_n)}{f''(\theta_n)}$$

Direct Numeric Optimization III

- If θ is multivariate, then again we use the gradient: $\nabla f(\theta)$, which is a vector, and Hessian: $\nabla^2 f(\theta)$, which is a matrix:

$$\theta_{n+1} = \theta_n - \left(\nabla^2 f(\theta_n)\right)^{-1} \nabla\left(f(\theta_n)\right).$$

- This is the basic approach of many traditional machine learning algorithms:
  - Pick a model that depends on θ, and a loss-function $f(\theta)$ that depends on the data and model (here, our loss is the log-likelihood function).
  - If possible, calculate the derivative of the loss function with respect to θ. If not, approximate numerically.
  - If possible, calculate the Hessian of the loss function with respect to θ. If not, approximate numerically.

Direct Numeric Optimization IV

- - Start with initial guess for θ, take steps the size of the (approximate) hessian of $f(\theta)$, in the direction of the gradient $f(\theta)$.

- Many optimization strategies are variants of this basic approach: In practice, Hessians / Gradients may be expensive to compute.

- We will primarily focus our attention on pre-existing solutions.

- In R, the function we will use is called optim.

- There are several optimization methods available within this function

Direct Numeric Optimization V

- Nelder-Mead: a heuristic, direct search method. Does not require differentiability. Slow, but robust.
- BFGS / L-BFGS-B: A Quasi-Newton approach, approximates either (or both) the Hessian and Gradient numerically. Extremely effective for (twice) differentiable functions, but slow as θ grows in dimension.
- SANN: A stochastic approach. Hard to tune, but effective on "rough" surfaces
- CG: Conjugate Gradients. More "fragile" than BFGS, but require less memory storage so can be useful for large θ.
- Brent: Only univariate θ. In these cases, approximates the gradient and hessian, similar to what we proposed. Very effective, but requires univariate θ.

Direct Numeric Optimization VI

- How it works: Get function f, method you want to use, and starting point θ_0. Example:

```r
gamma_negloglik <- function(theta) {
  -dgamma(
    x = data,
    shape = theta[1], rate = theta[2],
    log = TRUE
  ) |> sum()
}
```

```
optim(
  c(2, 5),
  fn = gamma_negloglik,
  method = 'L-BFGS-B',  # Constrained theta>0
  lower = c(1e-8, 1e-8)  # Constrained theta>0
)

$par
[1] 0.9728026 1.9806150

$value
[1] 6.641716

$counts
function gradient
      19       19
```

Gradient Descent and Variations

- In Newton-Raphson, the Hessian just tells us the *size* of the step.
- Often the gradient is available, but not the Hessian (or it is expensive to compute the inverse, since it is a matrix).
- In these cases, we often use a different approach called gradient descent.
- Idea: still step in the direction of the (negative) gradient, but the step size will just be small δ_n, and let δ_n shrink over-time:

$$\theta_{n+1} = \theta_n - \delta_n \nabla f(\theta_n).$$

Gradient Descent and Variations II

- The idea is simple, but a lot of modern optimization techniques and theory are based on this idea, finding various strategies for specifying δ_n.

- Ex: BFGS approximates the Hessian, but scales poorly with dimension. In deep learning, θ has dimension in millions / billions, so BFGS won't work (nor does `optim`).

- The success of deep-learning and AI is largely due to the advent of automatic differentiation: software that calculates the exact gradient of $f(\theta_n)$, after simple calculations.

- Auto-diff libraries: PyTorch, TensorFlow, JAX, etc. Mostly available in Python, though some are available in R.

Gradient Descent and Variations III

- Final approach we will discuss is stochastic gradient descent, which is effectively randomly sampling the data to compute an approximate gradient.

- The idea is that, even with auto-diff, gradients (with entire data) are expensive to compute.

- Thus, randomly sample data to get a stochastic approximation of the gradient; this is faster to compute, so we get more update-steps.

- This is the primary technique used in most deep-learning frameworks.

Maximum Likelihood: continued examples

Muon Decay

Returning to the Muon-Decay example we used MoM to solve.

The density of iid observations $X_1, \ldots, X_n$ is:

$$f(x; \alpha) = \frac{1 + \alpha x}{2}, \quad -1 \leq x \leq 1 \quad -1 \leq \alpha \leq 1.$$

Find the MLE of α.

Maximum Likelihood: continued examples II

"If [the likelihood] cannot by maximized analytically, it may be possible to use a computer to maximize [the likelihood] numerically. In fact, this is one of the most important features of the MLE. If a model can be written down, then there is some hope of maximizing it numerically, and, hence, finding MLEs of the parameters." – *Casella and Berger (2024, Section 7.2)*

Constraint Optimization

Constrained optimization

- In all of our previous examples, we want to maximize the log-likelihood function $f(x_i^*; \theta)$, and there is some natural constraints of θ.

- Examples: In Gamma(α, λ), $\theta = (\alpha, \lambda)$, and we are constrained by $\alpha, \lambda > 0$.

- In the Gaussian model, $\theta = (\mu, \sigma)$. μ is unconstrained, but $\sigma \geq 0$.

- Similarly, in the Muon-decay example: $\theta = \alpha \in (0, 1)$.

- This leads to the issue of *constrained* optimization, and is typically a requirement for maximum likelihood estimation.

- There are a few common strategies:

Constrained optimization II

- Limit search region to be within constraints: This is the basic approach used by the uniroot function. It's also used sometimes in optim, e.g., the L-BFGS-B method.
- Variable transformations. Modify the function so that the search algorithm can try all possible values, but apply a transformation that keeps it in the desired range.
    - Example: if $\theta > 0$ is a constraint, optimize over α with $\theta = e^\alpha$. Thus, $\alpha < 0$ is no problem, since $\theta > 0$. (Example on next slide)
    - Another common example is a logit transformation, which ensures $0 \leq \theta \leq 1$, via $\log(\theta/(1-\theta)) = \alpha$.
- In some cases, the constraint leads to an equation $g(\theta) = 0$, in which case we can apply something like Lagrange-multipliers. We will look at one of these cases as well.

Example: Gamma likelihood

MLE of Gamma likelihood using transformations

Revisit the numeric optimization of the Gamma likelihood, using variable transformations to deal with constraints.

```
gamma_negloglik2 <- function(alpha) {
  theta <- exp(alpha)
  -dgamma(
    x = data,
    shape = theta[1], rate = theta[2],
    log = TRUE
  ) |> sum()
}
```

Example: Gamma likelihood II

```
results <- optim(
  c(log(2), log(5)),
  fn = gamma_negloglik2,
  method = 'BFGS' # unconstrained
)
exp(results$par) # Theta

[1] 0.9728029 1.9806131
```

Multinomial Cell probabilities

- In this example, we are bound by a curve, not a region. Lagrange Multipliers are a good analytic approach.

- Consider fitting a multinomial distribution to a frequency table.

- We let $X_1, \ldots, X_m$ be the counts in cells $1, \ldots, X_m$, and we assume that $(X_1, \ldots, X_m)$ follows a multinomial distribution.

- The distribution has has parameters n: total counts, and p_i probability of cell i.

- The total count $n = \sum_{i=1}^{b} X_i$, and p_i being the

- In this case, the X_i are *not* independent, so we don't just take product of IID observations.

Multinomial Cell probabilities II

- However, the likelihood can be calculated via the pmf of the multinomial distribution:

$$L(\theta) = f(x_i^*; \theta) = \frac{n!}{\prod_{i=1}^{m} x_i!} \prod_{i=1}^{m} p_i^{x_i^*}.$$

- Given $X_1, \ldots, X_m$, the number of trials n is known. The parameter vector of interest is then:

$$\theta = (p_1, \ldots, p_m).$$

- The log-likelihood is given by:

$$\ell(\theta) = \log n! - \sum_{i=1}^{m} \log x_i! + \sum_{i=1}^{m} x_i \log p_i.$$

Multinomial Cell probabilities III

- Note that we have the constraint $\sum_i p_i = 1$. Thus, we will solve this using the Lagrange-multiplier technique.

- Reminder: Suppose we are trying to maximize (or minimize) a function $f(\theta)$, with the constraint $g(\theta) = 0$. That is,

$$\max_\theta f(\theta)$$
$$\text{subject to } g(\theta) = 0$$

- This optimization problem can be solved by introducing the Lagrange function:

$$\mathcal{L}(\theta, \lambda) = f(\theta) + \lambda g(\theta),$$

Multinomial Cell probabilities IV

- And solving the system of equations that arises from:

$$\nabla_{\theta, \lambda} \mathcal{L}(\theta, \lambda) = 0.$$

- Recall the partial derivative $\frac{\partial}{\partial \lambda} \mathcal{L}(\theta, \lambda) = g(\theta)$, which must be zero, giving the necessary constraint.

- Then, the condition $\nabla_{\theta} \mathcal{L}(\theta, \lambda) = 0$ ensures that the gradients of f and g are parallel. That is, geometrically, the constraint $g(\theta) = 0$ defines a curve or surface that we are constrained to; on this surface, the max/min value of f occurs when a level-set of f is tangent to the constraint curve $g(\theta) = 0$. This occurs only where $\nabla f(\theta) = \lambda \nabla g(\theta)$, which gives rise to the set of equations defined by $\nabla_{\theta} \mathcal{L}(\theta, \lambda) = 0$.

MLE of multinomial cell probabilities

Use Lagrange multipliers to find the MLE of a multinomial distribution, with $X_1, \ldots, X_m$.

Multinomial Cell probabilities VI

- Often, we can describe parameters of a model as functions of other parameters.

- Example: in the multinomial model, we have parameters p_i representing probabilities for each cell. We might want to make these a function of θ, such that $p_i(\theta)$; then, we will still want to estimate θ using the data.

- In the case above the log-likelihood is then:

$$\ell(\theta) = \log n! - \sum_{i=1}^{m} \log x_i^*! + \sum i = 1^m x_i \log p_i(\theta).$$

Hardy-Weinberg Equilibrium

- Consider alleles expressed as dominant (A) and recessive (a), and they are expressed with rates $1 - \theta$ and θ, respectively.

- The Hardy-Weinberg principle states genotypes are expressed in large, randomly mixing populations following the Punnett square in Table 1.

	A	a
A	$(1-p)^2$	$p(1-p)$
a	$p(1-p)$	p^2

Table 1: Genotype frequencies in large, randomly mixing population.

Hardy-Weinberg Equilibrium II

- Since the order of Aa or aA doesn't matter, this results in frequencies for AA, Aa, and aa: $(1-\theta)^2, 2\theta(1-\theta), \theta^2$.

- Supposing that we observe a population with these three possible gene expressions. Then, using population counts, we want to fit a multinomial model with $m = 3$ categories, and probabilities defined as:

$$p_1 = (1-\theta)^2, \quad p_2 = 2\theta(1-\theta), \quad p_3 = \theta^2.$$

Blood Types

In a Chinese population of Hong Kong in 1937, blood types occurred with the following frequencies, where M and N are erythrocyte antigens: M (342), MN (500), N (187), for a total of 1029 samples. Use a multinomial model to estimate the frequency of alleles and genotypes.

Invariance property of the MLE

Invariance property

- We will next talk about general properties of estimators, and introduce some theory.

- Before we do that, we still want to discuss point-estimation following the Bayesian approach.

- The order of things is a little tricky, but before moving on we will introduce one property of the MLE that can be used in finding the MLE.

- The property is known as invariance.

- Idea: information we have about θ should be the same information we have about $g(\theta)$ – we don't gain or loose information by transforming the variable.

Invariance property II

Theorem: Invariance property of the MLE

(Theorem 7.2.10, Casella and Berger, 2024) If $\hat{\theta}$ is the MLE of θ, then for any function g, $g(\hat{\theta})$ is the MLE of $g(\theta)$. For instance, if $\hat{\theta} = \bar{x}_n^*$, then the MLE of $\sin(\theta)$ is $\sin(\bar{x}_n^*)$.

Invariance property III

- Note in the previous theorem, nothing is preventing θ to be multivariate.

- Thus, if $\hat{\theta} = (\hat{\theta}_1, \ldots, \hat{\theta}_k)$ is the MLE of θ, then $g(\hat{\theta}) = \left(g(\hat{\theta}_1), \ldots, g(\hat{\theta}_k) \right)$ is the MLE of $g(\theta)$.

- For a more concrete example: suppose that p is the probability of an event, a model parameter that we want to estimate. Then if $\hat{p}$ is the MLE, then we can find the MLE for the odds-ratio or log-odds ratio very easily:

$$\text{MLE of odds ratio} = \frac{\hat{p}}{1 - \hat{p}}, \quad \text{MLE of log-odds ratio} = \log \frac{\hat{p}}{1 - \hat{p}}.$$

Comments on maximum likelihood

- There are many advantages of MLE, and is one reason that MLE is one of the gold-standards of statistical estimation (the other dominant approach being Bayesian estimation).

- Maximum likelihood is generally considered a frequentist approach due to asymptotic theory results, which we cover later. As presented, however, the MLE has more of a Bayesian flavor to it: the likelihood measures some degree of "belief" about a parameter, and we estimate the parameter by picking the value that maximizes this belief.

- Maximum likelihood is not without it's criticisms. For instance:

Comments on maximum likelihood II

- The MLE may not be unique (i.e., more than one global maximum). For instance, $\hat{\theta}_1$ and $\hat{\theta}_2$ may both maximize the likelihood $L(\theta)$; if they are very different estimates, this could cause a problem. It's also possible for an infinite number of θ values to maximize the likelihood, which again can be problematic.
- The MLE is often biased, and this can be particularly problematic in small sample sizes.
- The MLE depends on model assumptions. If they are wrong, the estimates might not make much sense.
- The MLE might not even exist! Le Cam, a famous 20th century statistician who was largely a proponent of the MLE approach, wrote a fun paper about some of these weaknesses, and warned against blindly using the MLE (Le Cam, 1990):

Comments on maximum likelihood III

"We present a list of principles leading to the construction of good estimates. The main principle says that one should not believe in principles but study each problem for its own sake." " – *(Le Cam, 1990)*

- Even if likelihoods can be maximized numerically, care is needed to ensure that final estimates correspond to global maximums and not just local solutions.

Comments on maximum likelihood IV

- Other difficulties arise in more complex models, such as time-series or spatial data. Here, the data are not IID, and the distribution of each observation changes over times / space, and observations are not independent. As a result, it's generally very difficult to write down the likelihood, or find the MLE. Minimizing a loss-function is generally easier in complex settings, even if it results in a loss of information (discussed later), or is equivalent to unrealistic Gaussian approximations (i.e., minimizing MSE).

References and Acknowledgements

Casella G, Berger R (2024). *Statistical inference*. Chapman and Hall/CRC.

Le Cam L (1990). "Maximum likelihood: an introduction." *International Statistical Review/Revue Internationale de Statistique*, pp. 153–171.

Pawitan Y (2001). *In all likelihood: statistical modelling and inference using likelihood*. Oxford University Press.

Rice JA (2007). *Mathematical statistics and data analysis*, volume 371. 3 edition. Thomson/Brooks/Cole Belmont, CA.

References and Acknowledgements II

- Compiled on January 27, 2026 using R version 4.5.2.

- We acknowledge students and instructors for previous versions of this course / slides.